

# Machine Learning with PyTorch and Scikit-Learn

## -- Code Examples

### Package version checks

Add folder to path in order to load from the check\_packages.py script:

```
In [1]: import sys
        sys.path.insert(0, '..')
```

Check recommended package versions:

```
In [4]: from python_environment_check import check_packages

        d = {
            'numpy': '1.21.2',
            'matplotlib': '3.4.3',
            'sklearn': '1.0',
            'pandas': '1.3.2'
        }
        check_packages(d)
```

## Chapter 3 - A Tour of Machine Learning Classifiers Using Scikit-Learn

### Overview

- [Choosing a classification algorithm](#)
- [First steps with scikit-learn](#)
  - [Training a perceptron via scikit-learn](#)
- [Modeling class probabilities via logistic regression](#)
  - [Logistic regression intuition and conditional probabilities](#)
  - [Learning the weights of the logistic loss function](#)
  - [Training a logistic regression model with scikit-learn](#)
  - [Tackling overfitting via regularization](#)
- [Maximum margin classification with support vector machines](#)
  - [Maximum margin intuition](#)
  - [Dealing with the nonlinearly separable case using slack variables](#)
  - [Alternative implementations in scikit-learn](#)
- [Solving nonlinear problems using a kernel SVM](#)
  - [Using the kernel trick to find separating hyperplanes in higher dimensional space](#)

- [Decision tree learning](#)
  - [Maximizing information gain – getting the most bang for the buck](#)
  - [Building a decision tree](#)
  - [Combining weak to strong learners via random forests](#)
- [K-nearest neighbors – a lazy learning algorithm](#)
- [Summary](#)

```
In [3]: from IPython.display import Image
        %matplotlib inline
```

## Choosing a classification algorithm

...

## First steps with scikit-learn

Loading the Iris dataset from scikit-learn. Here, the third column represents the petal length, and the fourth column the petal width of the flower examples. The classes are already converted to integer labels where 0=Iris-Setosa, 1=Iris-Versicolor, 2=Iris-Virginica.

```
In [4]: from sklearn import datasets
        import numpy as np

        iris = datasets.load_iris()
        X = iris.data[:, [2, 3]]
        y = iris.target

        print('Class labels:', np.unique(y))
```

Class labels: [0 1 2]

Splitting data into 70% training and 30% test data:

```
In [5]: from sklearn.model_selection import train_test_split

        X_train, X_test, y_train, y_test = train_test_split(
            X, y, test_size=0.3, random_state=1, stratify=y)
```

```
In [6]: print('Labels counts in y:', np.bincount(y))
        print('Labels counts in y_train:', np.bincount(y_train))
        print('Labels counts in y_test:', np.bincount(y_test))
```

Labels counts in y: [50 50 50]  
Labels counts in y\_train: [35 35 35]  
Labels counts in y\_test: [15 15 15]

Standardizing the features:

```
In [7]: from sklearn.preprocessing import StandardScaler

sc = StandardScaler()
sc.fit(X_train)
X_train_std = sc.transform(X_train)
X_test_std = sc.transform(X_test)
```

## Training a perceptron via scikit-learn

```
In [9]: from sklearn.linear_model import Perceptron

ppn = Perceptron(eta0=0.1, random_state=1)
ppn.fit(X_train_std, y_train)
```

```
Out[9]: Perceptron(eta0=0.1, random_state=1)
```

```
In [10]: y_pred = ppn.predict(X_test_std)
print('Misclassified examples: %d' % (y_test != y_pred).sum())

Misclassified examples: 1
```

```
In [11]: from sklearn.metrics import accuracy_score

print('Accuracy: %.3f' % accuracy_score(y_test, y_pred))

Accuracy: 0.978
```

```
In [12]: print('Accuracy: %.3f' % ppn.score(X_test_std, y_test))

Accuracy: 0.978
```

```
In [13]: from matplotlib.colors import ListedColormap
import matplotlib.pyplot as plt

# To check recent matplotlib compatibility
import matplotlib
from distutils.version import LooseVersion

def plot_decision_regions(X, y, classifier, test_idx=None, resolution=0.02):

    # setup marker generator and color map
    markers = ('o', 's', '^', 'v', '<')
    colors = ('red', 'blue', 'lightgreen', 'gray', 'cyan')
    cmap = ListedColormap(colors[:len(np.unique(y))])

    # plot the decision surface
    x1_min, x1_max = X[:, 0].min() - 1, X[:, 0].max() + 1
    x2_min, x2_max = X[:, 1].min() - 1, X[:, 1].max() + 1
    xx1, xx2 = np.meshgrid(np.arange(x1_min, x1_max, resolution),
                           np.arange(x2_min, x2_max, resolution))
    lab = classifier.predict(np.array([xx1.ravel(), xx2.ravel()]).T)
    lab = lab.reshape(xx1.shape)
    plt.contourf(xx1, xx2, lab, alpha=0.3, cmap=cmap)
    plt.xlim(xx1.min(), xx1.max())
    plt.ylim(xx2.min(), xx2.max())
```

```

# plot class examples
for idx, cl in enumerate(np.unique(y)):
    plt.scatter(x=X[y == cl, 0],
                y=X[y == cl, 1],
                alpha=0.8,
                c=colors[idx],
                marker=markers[idx],
                label=f'Class {cl}',
                edgecolor='black')

# highlight test examples
if test_idx:
    # plot all examples
    X_test, y_test = X[test_idx, :], y[test_idx]

    plt.scatter(X_test[:, 0],
                X_test[:, 1],
                c='none',
                edgecolor='black',
                alpha=1.0,
                linewidth=1,
                marker='o',
                s=100,
                label='Test set')

```

Training a perceptron model using the standardized training data:


```

In [15]: X_combined_std = np.vstack((X_train_std, X_test_std))
         y_combined = np.hstack((y_train, y_test))

         plot_decision_regions(X=X_combined_std, y=y_combined,
                               classifier=ppn, test_idx=range(105, 150))
         plt.xlabel('Petal length [standardized]')
         plt.ylabel('Petal width [standardized]')
         plt.legend(loc='upper left')

         plt.tight_layout()
         #plt.savefig('figures/03_01.png', dpi=300)
         plt.show()

```



# Modeling class probabilities via logistic regression

...

## Logistic regression intuition and conditional probabilities

```
In [17]: import matplotlib.pyplot as plt
import numpy as np


def sigmoid(z):
    return 1.0 / (1.0 + np.exp(-z))

z = np.arange(-7, 7, 0.1)
sigma_z = sigmoid(z)

plt.plot(z, sigma_z)
plt.axvline(0.0, color='k')
plt.ylim(-0.1, 1.1)
plt.xlabel('z')
plt.ylabel('$\sigma(z)$')

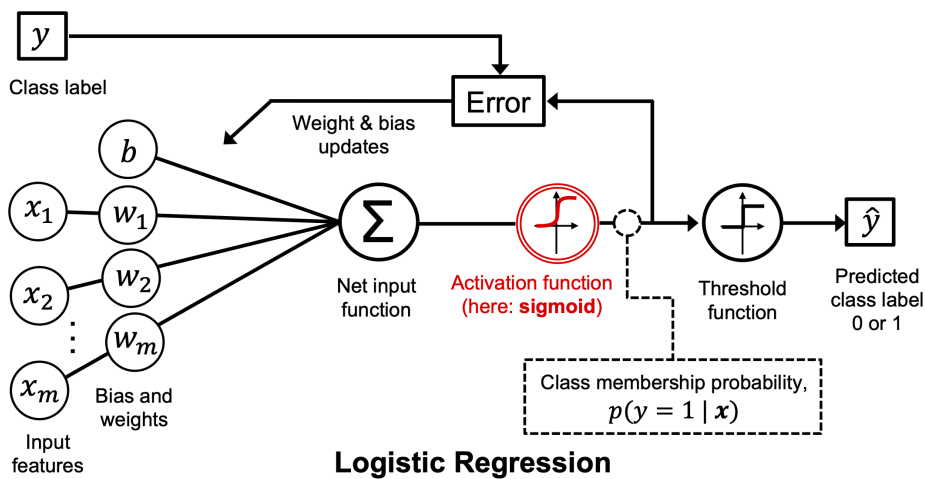
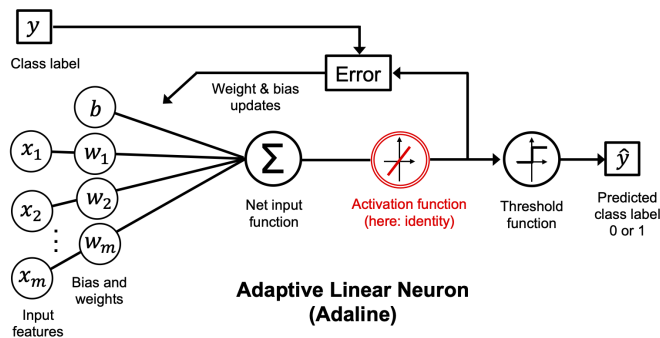
# y axis ticks and gridline
plt.yticks([0.0, 0.5, 1.0])
ax = plt.gca()
ax.yaxis.grid(True)

plt.tight_layout()
# plt.savefig('figures/03_02.png', dpi=300)
plt.show()
```



```
In [18]: Image(filename='figures/03_03.png', width=500)
```

Out[18]:

In [19]: `Image(filename='figures/03_25.png', width=500)`

Out[19]:

$$\frac{\partial L}{\partial w_j} = \underbrace{\frac{\partial L}{\partial a} \frac{da}{dz} \frac{\partial z}{\partial w_j}}_{\text{Apply chain rule}} \quad \text{where } a = \sigma(z) = \frac{1}{1 + e^{-z}}$$

1) Derive terms separately:

2) Combine via chain rule and simplify:

$$\begin{aligned} \frac{\partial L}{\partial a} &= \frac{a - y}{a - a^2} \\ \frac{da}{dz} &= \frac{e^{-z}}{(1 + e^{-z})^2} = a \cdot (1 - a) \\ \frac{\partial z}{\partial w_j} &= x_j \end{aligned}$$

$$\frac{\partial L}{\partial z} = a - y$$

$$\frac{\partial L}{\partial w_j} = (a - y)x_j = -(y - a)x_j$$

## Learning the weights of the logistic loss function

```
In [57]: def loss_1(z):
          return - np.log(sigmoid(z))

          def loss_0(z):
```

```

    return - np.log(1 - sigmoid(z))


z = np.arange(-10, 10, 0.1)
sigma_z = sigmoid(z)

c1 = [loss_1(x) for x in z]
plt.plot(sigma_z, c1, label='L(w, b) if y=1')

c0 = [loss_0(x) for x in z]
plt.plot(sigma_z, c0, linestyle='--', label='L(w, b) if y=0')

plt.ylim(0.0, 5.1)
plt.xlim([0, 1])
plt.xlabel('$\sigma(z)$')
plt.ylabel('L(w, b)')
plt.legend(loc='best')
plt.tight_layout()
#plt.savefig('figures/03_04.png', dpi=300)
plt.show()

```



```

In [22]: class LogisticRegressionGD:
    """Gradient descent-based logistic regression classifier.

    Parameters
    -----
    eta : float
        Learning rate (between 0.0 and 1.0)
    n_iter : int
        Passes over the training dataset.
    random_state : int
        Random number generator seed for random weight
        initialization.

    Attributes
    -----
    w_ : 1d-array
        Weights after training.
    b_ : Scalar
        Bias unit after fitting.
    losses_ : list
        Log loss function values in each epoch.

    """
    def __init__(self, eta=0.01, n_iter=50, random_state=1):
        self.eta = eta

```

```

self.n_iter = n_iter
self.random_state = random_state

def fit(self, X, y):
    """ Fit training data.

    Parameters
    -----
    X : {array-like}, shape = [n_examples, n_features]
        Training vectors, where n_examples is the number of examples and
        n_features is the number of features.
    y : array-like, shape = [n_examples]
        Target values.

    Returns
    -----
    self : Instance of LogisticRegressionGD

    """
    rgen = np.random.RandomState(self.random_state)
    self.w_ = rgen.normal(loc=0.0, scale=0.01, size=X.shape[1])
    self.b_ = np.float_(0.)
    self.losses_ = []

    for i in range(self.n_iter):
        net_input = self.net_input(X)
        output = self.activation(net_input)
        errors = (y - output)
        self.w_ += self.eta * X.T.dot(errors) / X.shape[0]
        self.b_ += self.eta * errors.mean()
        loss = (-y.dot(np.log(output)) - (1 - y).dot(np.log(1 - output))) / X.shape[0]
        self.losses_.append(loss)
    return self

def net_input(self, X):
    """Calculate net input"""
    return np.dot(X, self.w_) + self.b_

def activation(self, z):
    """Compute logistic sigmoid activation"""
    return 1. / (1. + np.exp(-np.clip(z, -250, 250)))

def predict(self, X):
    """Return class label after unit step"""
    return np.where(self.activation(self.net_input(X)) >= 0.5, 1, 0)

```

```

In [24]: X_train_01_subset = X_train_std[(y_train == 0) | (y_train == 1)]
y_train_01_subset = y_train[(y_train == 0) | (y_train == 1)]

lrgd = LogisticRegressionGD(eta=0.3, n_iter=1000, random_state=1)
lrgd.fit(X_train_01_subset,
        y_train_01_subset)


plot_decision_regions(X=X_train_01_subset,
                    y=y_train_01_subset,
                    classifier=lrgd)

plt.xlabel('Petal length [standardized]')
plt.ylabel('Petal width [standardized]')
plt.legend(loc='upper left')

```



```
plt.tight_layout()
#plt.savefig('figures/03_05.png', dpi=300)
plt.show()
```




## Training a logistic regression model with scikit-learn

```
In [26]: from sklearn.linear_model import LogisticRegression

lr = LogisticRegression(C=100.0, solver='lbfgs', multi_class='ovr')
lr.fit(X_train_std, y_train)

plot_decision_regions(X_combined_std, y_combined,
                      classifier=lr, test_idx=range(105, 150))
plt.xlabel('Petal length [standardized]')
plt.ylabel('Petal width [standardized]')
plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('figures/03_06.png', dpi=300)
plt.show()
```



```
In [27]: lr.predict_proba(X_test_std[:3, :])
```

```
Out[27]: array([[3.81527885e-09, 1.44792866e-01, 8.55207131e-01],
 [8.34020679e-01, 1.65979321e-01, 3.25737138e-13],
 [8.48831425e-01, 1.51168575e-01, 2.62277619e-14]])
```

```
In [28]: lr.predict_proba(X_test_std[:3, :]).sum(axis=1)
```

```
Out[28]: array([1., 1., 1.])
```

```
In [29]: lr.predict_proba(X_test_std[:3, :]).argmax(axis=1)
```

```
Out[29]: array([2, 0, 0])
```

```
In [30]: lr.predict(X_test_std[:3, :])
```

```
Out[30]: array([2, 0, 0])
```


```
In [31]: lr.predict(X_test_std[0, :].reshape(1, -1))
```

```
Out[31]: array([2])
```

## Tackling overfitting via regularization


```
In [33]: Image(filename='figures/03_07.png', width=700)
```

```
Out[33]:
```



```
In [34]: weights, params = [], []
for c in np.arange(-5, 5):
    lr = LogisticRegression(C=10.**c,
                           multi_class='ovr')
    lr.fit(X_train_std, y_train)
    weights.append(lr.coef_[1])
    params.append(10.**c)


weights = np.array(weights)
plt.plot(params, weights[:, 0],
         label='Petal length')
plt.plot(params, weights[:, 1], linestyle='--',
         label='Petal width')
plt.ylabel('Weight coefficient')
plt.xlabel('C')
plt.legend(loc='upper left')
plt.xscale('log')
#plt.savefig('figures/03_08.png', dpi=300)
plt.show()
```



## Maximum margin classification with support vector machines

In [36]: `Image(filename='figures/03_09.png', width=700)`

Out[36]:



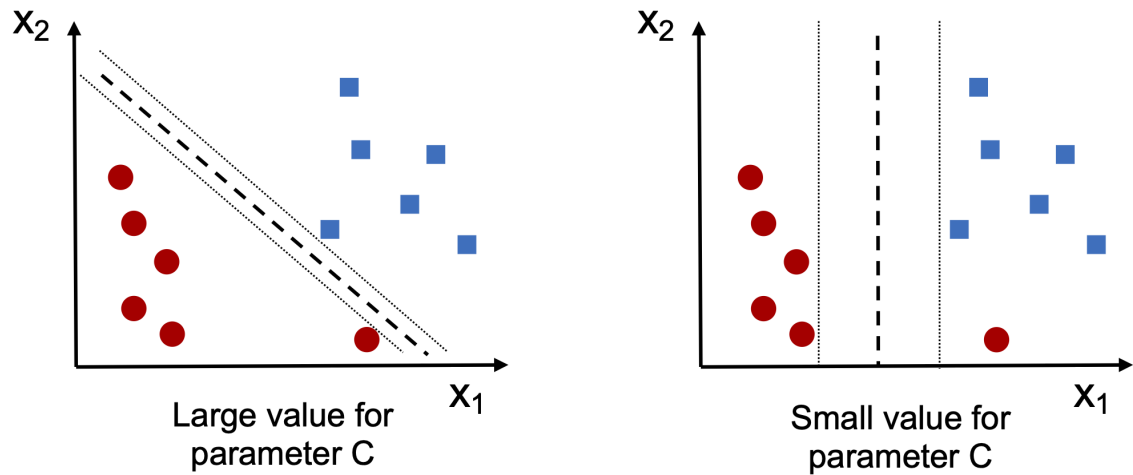
## Maximum margin intuition

...

## Dealing with the nonlinearly separable case using slack variables

In [38]: `Image(filename='figures/03_10.png', width=600)`

Out[38]:




```
In [39]: from sklearn.svm import SVC

svm = SVC(kernel='linear', C=1.0, random_state=1)
svm.fit(X_train_std, y_train)

plot_decision_regions(X_combined_std,
                      y_combined,
                      classifier=svm,
                      test_idx=range(105, 150))

plt.xlabel('Petal length [standardized]')
plt.ylabel('Petal width [standardized]')
plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('figures/03_11.png', dpi=300)
plt.show()
```



## Alternative implementations in scikit-learn

```
In [40]: from sklearn.linear_model import SGDClassifier

ppn = SGDClassifier(loss='perceptron')
lr = SGDClassifier(loss='log')
svm = SGDClassifier(loss='hinge')
```

# Solving non-linear problems using a kernel SVM


```
In [41]: import matplotlib.pyplot as plt
import numpy as np

np.random.seed(1)
X_xor = np.random.randn(200, 2)
y_xor = np.logical_xor(X_xor[:, 0] > 0,
                       X_xor[:, 1] > 0)
y_xor = np.where(y_xor, 1, 0)

plt.scatter(X_xor[y_xor == 1, 0],
            X_xor[y_xor == 1, 1],
            c='royalblue',
            marker='s',
            label='Class 1')
plt.scatter(X_xor[y_xor == 0, 0],
            X_xor[y_xor == 0, 1],
            c='tomato',
            marker='o',
            label='Class 0')

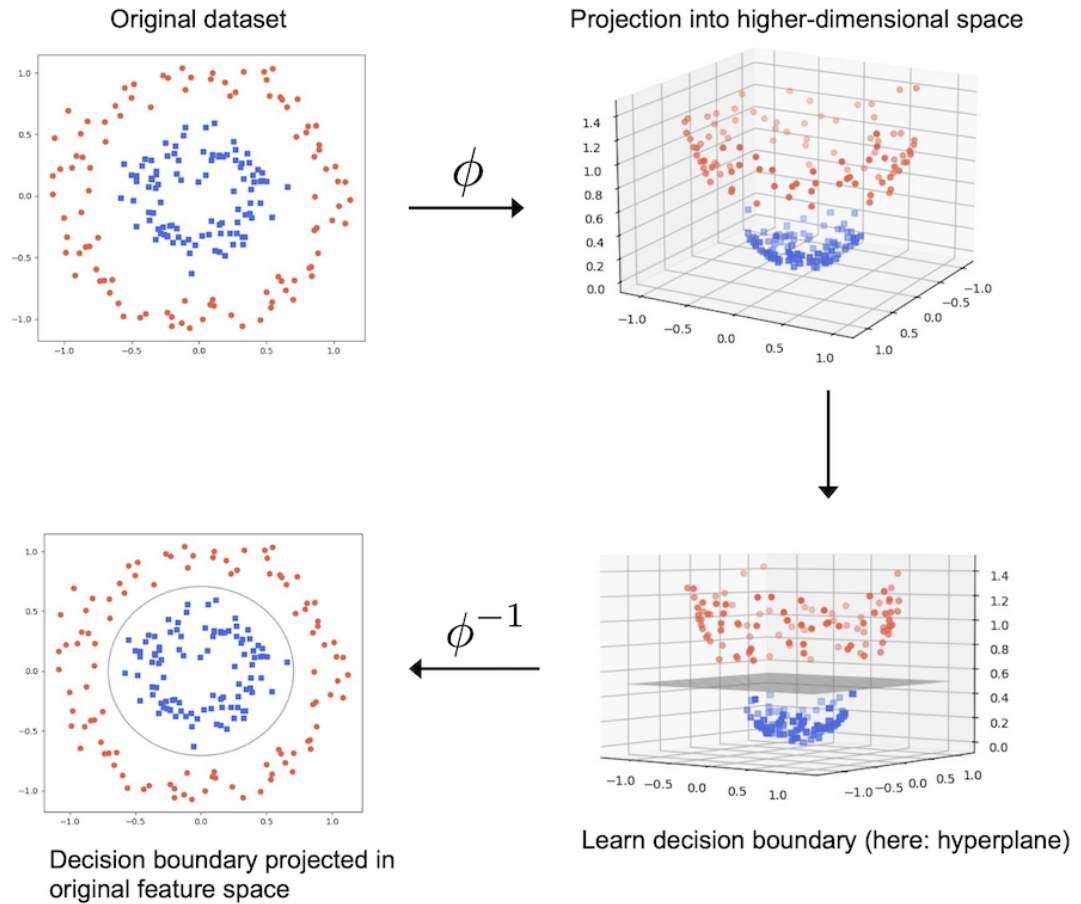
plt.xlim([-3, 3])
plt.ylim([-3, 3])
plt.xlabel('Feature 1')
plt.ylabel('Feature 2')

plt.legend(loc='best')
plt.tight_layout()
# plt.savefig('figures/03_12.png', dpi=300)
plt.show()
```



```
In [42]: Image(filename='figures/03_13.png', width=700)
```


Out[42]:



## Using the kernel trick to find separating hyperplanes in higher dimensional space

```
In [43]: svm = SVC(kernel='rbf', random_state=1, gamma=0.10, C=10.0)
svm.fit(X_xor, y_xor)
plot_decision_regions(X_xor, y_xor,
                      classifier=svm)


plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('figures/03_14.png', dpi=300)
plt.show()
```



```
In [44]: from sklearn.svm import SVC


svm = SVC(kernel='rbf', random_state=1, gamma=0.2, C=1.0)
svm.fit(X_train_std, y_train)

plot_decision_regions(X_combined_std, y_combined,
                      classifier=svm, test_idx=range(105, 150))
plt.xlabel('Petal length [standardized]')
plt.ylabel('Petal width [standardized]')
plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('figures/03_15.png', dpi=300)
plt.show()
```



```
In [45]: svm = SVC(kernel='rbf', random_state=1, gamma=100.0, C=1.0)
svm.fit(X_train_std, y_train)

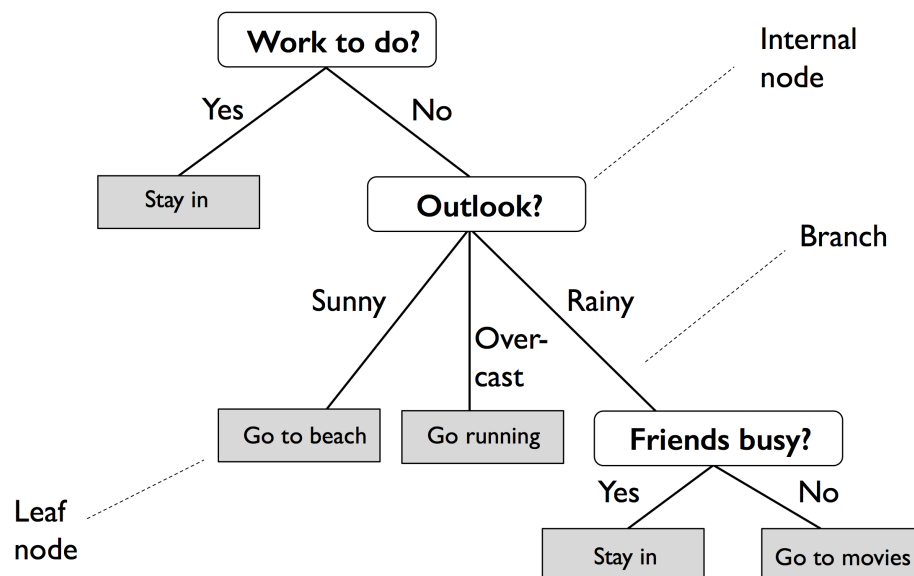
plot_decision_regions(X_combined_std, y_combined,
                      classifier=svm, test_idx=range(105, 150))
plt.xlabel('Petal length [standardized]')
plt.ylabel('Petal width [standardized]')
plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('figures/03_16.png', dpi=300)
plt.show()
```



## Decision tree learning

In [46]: `Image(filename='figures/03_17.png', width=500)`

Out[46]:




```
In [47]: def entropy(p):
    return - p * np.log2(p) - (1 - p) * np.log2((1 - p))

x = np.arange(0.0, 1.0, 0.01)
ent = [entropy(p) if p != 0 else None
        for p in x]

plt.ylabel('Entropy')
plt.xlabel('Class-membership probability p(i=1)')
plt.plot(x, ent)
#plt.savefig('figures/03_26.png', dpi=300)
plt.show()
```





In [48]: `Image(filename='figures/03_18.png', width=500)`

Out[48]:



## Maximizing information gain - getting the most bang for the buck

```

In [49]: import matplotlib.pyplot as plt
import numpy as np

def gini(p):
    return p * (1 - p) + (1 - p) * (1 - (1 - p))

def entropy(p):
    return - p * np.log2(p) - (1 - p) * np.log2((1 - p))

def error(p):
    return 1 - np.max([p, 1 - p])

x = np.arange(0.0, 1.0, 0.01)

ent = [entropy(p) if p != 0 else None for p in x]
sc_ent = [e * 0.5 if e else None for e in ent]
err = [error(i) for i in x]

fig = plt.figure()
ax = plt.subplot(111)
for i, lab, ls, c, in zip([ent, sc_ent, gini(x), err],
                        ['Entropy', 'Entropy (scaled)',
                        'Gini impurity', 'Misclassification error'],

```


```

        ['-', '--', '---', '-.-'],
        ['black', 'lightgray', 'red', 'green', 'cyan']):
    line = ax.plot(x, i, label=lab, linestyle=ls, lw=2, color=c)

ax.legend(loc='upper center', bbox_to_anchor=(0.5, 1.15),
          ncol=5, fancybox=True, shadow=False)

ax.axhline(y=0.5, linewidth=1, color='k', linestyle='--')
ax.axhline(y=1.0, linewidth=1, color='k', linestyle='--')
plt.ylim([0, 1.1])
plt.xlabel('p(i=1)')
plt.ylabel('Impurity index')
#plt.savefig('figures/03_19.png', dpi=300, bbox_inches='tight')
plt.show()

```



## Building a decision tree

```

In [50]: from sklearn.tree import DecisionTreeClassifier


tree_model = DecisionTreeClassifier(criterion='gini',
                                   max_depth=4,
                                   random_state=1)

tree_model.fit(X_train, y_train)

X_combined = np.vstack((X_train, X_test))
y_combined = np.hstack((y_train, y_test))
plot_decision_regions(X_combined, y_combined,
                     classifier=tree_model,
                     test_idx=range(105, 150))

plt.xlabel('Petal length [cm]')
plt.ylabel('Petal width [cm]')
plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('figures/03_20.png', dpi=300)
plt.show()

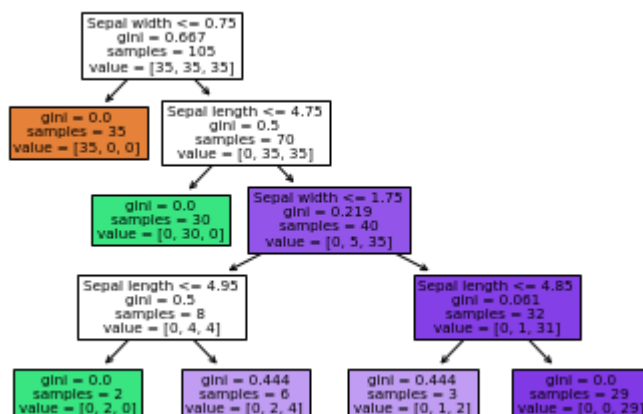
```



```
In [51]: from sklearn import tree

feature_names = ['Sepal length', 'Sepal width',
                 'Petal length', 'Petal width']
tree.plot_tree(tree_model,
               feature_names=feature_names,
               filled=True)

#plt.savefig('figures/03_21_1.pdf')
plt.show()
```



## Combining weak to strong learners via random forests

```
In [52]: from sklearn.ensemble import RandomForestClassifier

forest = RandomForestClassifier(n_estimators=25,
                              random_state=1,
                              n_jobs=2)


forest.fit(X_train, y_train)
```

```

plot_decision_regions(X_combined, y_combined,
                      classifier=forest, test_idx=range(105, 150))

plt.xlabel('Petal length [cm]')
plt.ylabel('Petal width [cm]')
plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('figures/03_2.png', dpi=300)
plt.show()


```



## K-nearest neighbors - a lazy learning algorithm

In [53]: `Image(filename='figures/03_23.png', width=400)`

Out[53]:




```
In [54]: from sklearn.neighbors import KNeighborsClassifier

knn = KNeighborsClassifier(n_neighbors=5,
                          p=2,
                          metric='minkowski')
knn.fit(X_train_std, y_train)

plot_decision_regions(X_combined_std, y_combined,
                      classifier=knn, test_idx=range(105, 150))

plt.xlabel('Petal length [standardized]')
plt.ylabel('Petal width [standardized]')
plt.legend(loc='upper left')
plt.tight_layout()
#plt.savefig('figures/03_24_figures.png', dpi=300)
plt.show()
```



## Summary

...

---

Readers may ignore the next cell.

```
In [55]: ! python ../convert_notebook_to_script.py --input ch03.ipynb --output ch03.py

[NbConvertApp] WARNING | Config option `kernel_spec_manager_class` not recognized
by `NbConvertApp`.
[NbConvertApp] Converting notebook ch03.ipynb to script
[NbConvertApp] Writing 19384 bytes to ch03.py
```

```
In [ ]:
```